

credit-score

July 17, 2024

```
[ ]: !pip install -q autoviz
      !pip install -q -U --pre pycaret
```

	67.5/67.5 kB
1.1 MB/s eta 0:00:00	431.4/431.4
kB 8.4 MB/s eta 0:00:00	2.0/2.0 MB
29.4 MB/s eta 0:00:00	255.9/255.9
MB 4.2 MB/s eta 0:00:00	155.4/155.4
kB 15.3 MB/s eta 0:00:00	24.7/24.7 MB
33.2 MB/s eta 0:00:00	8.3/8.3 MB
42.6 MB/s eta 0:00:00	7.0/7.0 MB
56.2 MB/s eta 0:00:00	486.1/486.1
kB 8.3 MB/s eta 0:00:00	302.2/302.2
kB 26.6 MB/s eta 0:00:00	13.4/13.4 MB
42.7 MB/s eta 0:00:00	163.8/163.8
kB 14.6 MB/s eta 0:00:00	
Preparing metadata (setup.py) ... done	258.3/258.3
kB 22.4 MB/s eta 0:00:00	
6.7 MB/s eta 0:00:00	81.9/81.9 kB

```

194.1/194.1
kB 12.9 MB/s eta 0:00:00
11.6/11.6 MB
85.1 MB/s eta 0:00:00
79.9/79.9 MB
9.7 MB/s eta 0:00:00
106.8/106.8
kB 11.0 MB/s eta 0:00:00
80.7/80.7 kB
8.5 MB/s eta 0:00:00
21.8/21.8 MB
56.9 MB/s eta 0:00:00
44.0/44.0 kB
4.5 MB/s eta 0:00:00
2.1/2.1 MB
69.9 MB/s eta 0:00:00
130.1/130.1
kB 12.6 MB/s eta 0:00:00
12.1/12.1 MB
89.1 MB/s eta 0:00:00
1.6/1.6 MB
72.4 MB/s eta 0:00:00
7.5/7.5 MB
98.0 MB/s eta 0:00:00
141.1/141.1
kB 13.0 MB/s eta 0:00:00
2.1/2.1 MB
81.5 MB/s eta 0:00:00

```

Building wheel for pyod (setup.py) ... done

```
[ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from pycaret import classification
from sklearn.model_selection import cross_val_score
from sklearn.metrics import classification_report, confusion_matrix
```

```
[17]: df = pd.read_csv('Credit Score Classification Dataset.csv')
```

```
[18]: df.head()
```

```
[18]:   Age  Gender  Income  Education  Marital Status \
0   25  Female   50000  Bachelor's Degree      Single
```

1	30	Male	100000	Master's Degree	Married
2	35	Female	75000	Doctorate	Married
3	40	Male	125000	High School Diploma	Single
4	45	Female	100000	Bachelor's Degree	Married

	Number of Children	Home Ownership	Credit Score
0	0	Rented	High
1	2	Owned	High
2	1	Owned	High
3	0	Owned	High
4	3	Owned	High

```
[19]: df.tail()
```

```
[19]:
```

	Age	Gender	Income	Education	Marital Status \
159	29	Female	27500	High School Diploma	Single
160	34	Male	47500	Associate's Degree	Single
161	39	Female	62500	Bachelor's Degree	Married
162	44	Male	87500	Master's Degree	Single
163	49	Female	77500	Doctorate	Married

	Number of Children	Home Ownership	Credit Score
159	0	Rented	Low
160	0	Rented	Average
161	2	Owned	High
162	0	Owned	High
163	1	Owned	High

## 1 1. Unique values

```
[20]: df.nunique()
```

```
[20]:
```

Age	29
Gender	2
Income	52
Education	5
Marital Status	2
Number of Children	4
Home Ownership	2
Credit Score	3

dtype: int64

## 2 2. Data Types

```
[21]: df.dtypes
```

```
[21]: Age                int64
      Gender            object
      Income            int64
      Education         object
      Marital Status    object
      Number of Children int64
      Home Ownership    object
      Credit Score      object
      dtype: object
```

```
[22]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 164 entries, 0 to 163
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Age                   164 non-null   int64
1   Gender                164 non-null   object
2   Income                164 non-null   int64
3   Education             164 non-null   object
4   Marital Status        164 non-null   object
5   Number of Children    164 non-null   int64
6   Home Ownership        164 non-null   object
7   Credit Score          164 non-null   object
dtypes: int64(3), object(5)
memory usage: 10.4+ KB
```

## 3 3. Statistics

```
[23]: df.describe().T
```

```
[23]:
```

	count	mean	std	min	25%	\
Age	164.0	37.975610	8.477289	25.0	30.75	
Income	164.0	83765.243902	32457.306728	25000.0	57500.00	
Number of Children	164.0	0.652439	0.883346	0.0	0.00	

	50%	75%	max
Age	37.0	45.0	53.0
Income	83750.0	105000.0	162500.0
Number of Children	0.0	1.0	3.0

## 4 4. Null values

```
[24]: df.isna().any()
```

```
[24]: Age                False
      Gender              False
      Income              False
      Education            False
      Marital Status      False
      Number of Children  False
      Home Ownership      False
      Credit Score        False
      dtype: bool
```

```
[25]: df.isna().sum()
```

```
[25]: Age                0
      Gender              0
      Income              0
      Education            0
      Marital Status      0
      Number of Children  0
      Home Ownership      0
      Credit Score        0
      dtype: int64
```

## 5 5. Categorical values

```
[26]: cat_cols = df.select_dtypes(include = ['object']).columns.tolist()
      cat_cols
```

```
[26]: ['Gender', 'Education', 'Marital Status', 'Home Ownership', 'Credit Score']
```

## 6 Label encoding Categorical values

```
[27]: from sklearn.preprocessing import LabelEncoder

      # Create a label encoder object
      le = LabelEncoder()

      for i in cat_cols:
          # Fit the label encoder object to the dataset
          le.fit(df[i])

          # Transform the dataset using the label encoder object
```

```
df[i] = le.transform(df[i])
# get list of categorical columns
cat_cols = df.select_dtypes(include=['object']).columns.tolist()
cat_cols
```

[27]: []

## 7 6. Correlation

[28]: df.corr()

[28]:

	Age	Gender	Income	Education	Marital Status \
Age	1.000000	0.235343	0.699464	0.170254	-0.517723
Gender	0.235343	1.000000	0.495738	0.248671	0.278362
Income	0.699464	0.495738	1.000000	0.369449	-0.471004
Education	0.170254	0.248671	0.369449	1.000000	-0.067797
Marital Status	-0.517723	0.278362	-0.471004	-0.067797	1.000000
Number of Children	0.055390	-0.442139	0.084547	0.047311	-0.696984
Home Ownership	-0.713803	-0.031519	-0.704928	-0.397043	0.708374
Credit Score	0.205362	-0.247729	0.083698	0.334424	-0.205756

	Number of Children	Home Ownership	Credit Score
Age	0.055390	-0.713803	0.205362
Gender	-0.442139	-0.031519	-0.247729
Income	0.084547	-0.704928	0.083698
Education	0.047311	-0.397043	0.334424
Marital Status	-0.696984	0.708374	-0.205756
Number of Children	1.000000	-0.497129	0.136517
Home Ownership	-0.497129	1.000000	-0.293384
Credit Score	0.136517	-0.293384	1.000000

[29]: # Set the size of figure to 12 by 10.  
plt.figure(figsize=(18,10))  
  
sns.heatmap(df.corr(), annot = True)

[29]: <Axes: >



## 8 7. Class Distributions

```
[30]: # Count the number of instances in each class
class_counts = df['Credit Score'].value_counts()

# Print the class distribution
print('Class distribution:')
print(class_counts)
```

```
Class distribution:
Credit Score
1      113
0       36
2       15
Name: count, dtype: int64
```

## 9 SMOTE

```
[31]: X = df.drop('Credit Score', axis = 1)
y = df['Credit Score']
from sklearn.model_selection import train_test_split

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳ random_state=42)
from imblearn.over_sampling import SMOTE
```

```

# Instantiate SMOTE
sm = SMOTE(random_state=42)

# Fit SMOTE to training data
X_train_res, y_train_res = sm.fit_resample(X_train, y_train)

# Print class distribution of original and resampled data
print('Class distribution before resampling:', y_train.value_counts())
print('Class distribution after resampling:', y_train_res.value_counts())

```

```

Class distribution before resampling: Credit Score
1    90
0    31
2    10
Name: count, dtype: int64
Class distribution after resampling: Credit Score
1    90
0    90
2    90
Name: count, dtype: int64

```

## 10 Other EDA Suggestions

```

[34]: from sklearn.impute import SimpleImputer
      from sklearn.preprocessing import StandardScaler
      def check_missing_values(df):
          missing_data = df.isnull().sum()
          missing_data = missing_data[missing_data > 0]
          missing_data.sort_values(ascending=False, inplace=True)
          print("Missing Values in Each Column:\n", missing_data)

          # Visualize missing data
          sns.heatmap(df.isnull(), cbar=False, cmap='viridis')
          plt.show()

      check_missing_values(df)
      def handle_missing_values(df):
          imputer = SimpleImputer(strategy='mean') # You can change strategy to
          ↪ 'median', 'most_frequent', or 'constant'
          df_imputed = pd.DataFrame(imputer.fit_transform(df), columns=df.columns)
          return df_imputed

      df = handle_missing_values(df)
      def check_duplicates(df):
          duplicates = df.duplicated().sum()

```



```

    print("Number of Duplicates: ", duplicates)
    return duplicates

duplicates = check_duplicates(df)
if duplicates > 0:
    df = df.drop_duplicates()
def scale_features(df):
    scaler = StandardScaler()
    numeric_columns = df.select_dtypes(include=[np.number]).columns
    df[numeric_columns] = scaler.fit_transform(df[numeric_columns])
    return df

df = scale_features(df)
def check_outliers(df):
    numeric_columns = df.select_dtypes(include=[np.number]).columns
    for col in numeric_columns:
        sns.boxplot(x=df[col])
        plt.title(f'Box plot for {col}')
        plt.show()

check_outliers(df)
def handle_outliers(df):
    # This example removes outliers using the IQR method
    numeric_columns = df.select_dtypes(include=[np.number]).columns
    for col in numeric_columns:
        Q1 = df[col].quantile(0.25)
        Q3 = df[col].quantile(0.75)
        IQR = Q3 - Q1
        df = df[~((df[col] < (Q1 - 1.5 * IQR)) | (df[col] > (Q3 + 1.5 * IQR)))]
    return df

df = handle_outliers(df)
def data_cleaning_suggestions(df):
    print("Step 1: Checking for Missing Values")
    check_missing_values(df)

    print("\nStep 2: Handling Missing Values")
    df = handle_missing_values(df)

    print("\nStep 3: Checking for Duplicates")
    duplicates = check_duplicates(df)

    if duplicates > 0:
        print("Removing Duplicates")
        df = df.drop_duplicates()

    print("\nStep 4: Scaling Numeric Features")

```

```

df = scale_features(df)

print("\nStep 5: Checking for Outliers")
check_outliers(df)

print("\nStep 6: Handling Outliers")
df = handle_outliers(df)


print("\nData Cleaning Completed")
return df

# Run the data cleaning suggestions
df_cleaned = data_cleaning_suggestions(df)

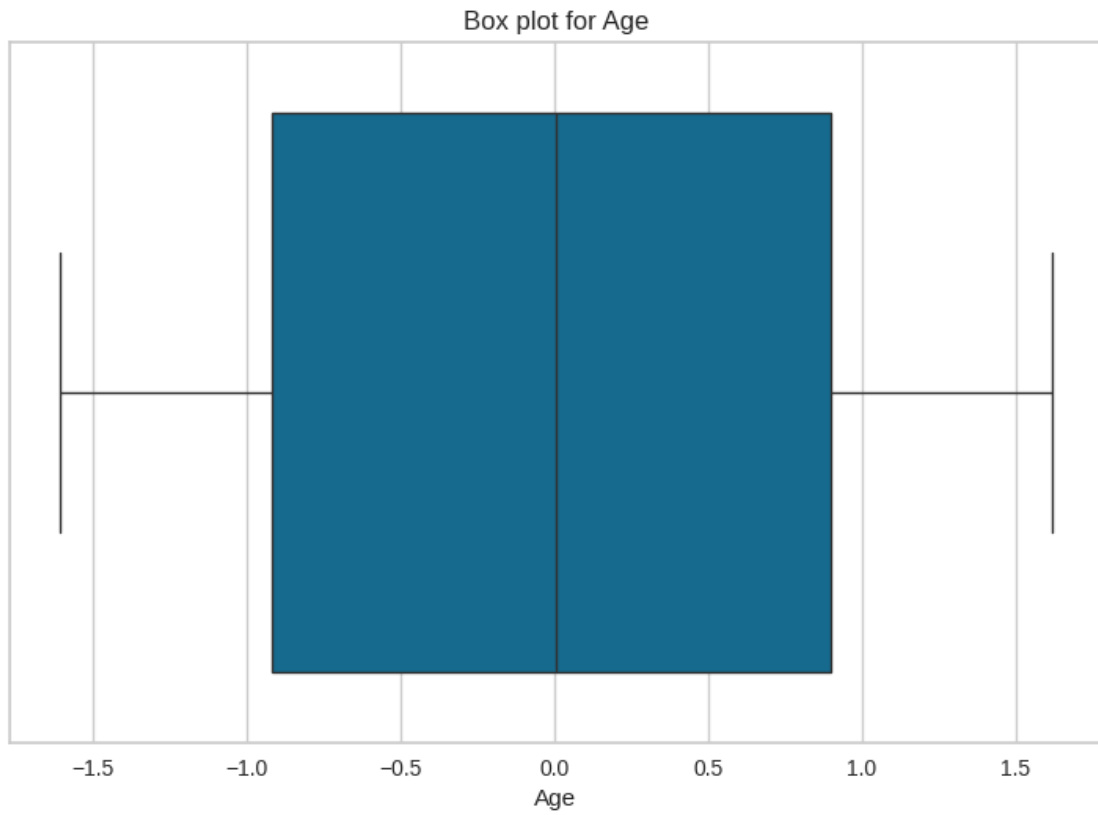
```


Missing Values in Each Column:

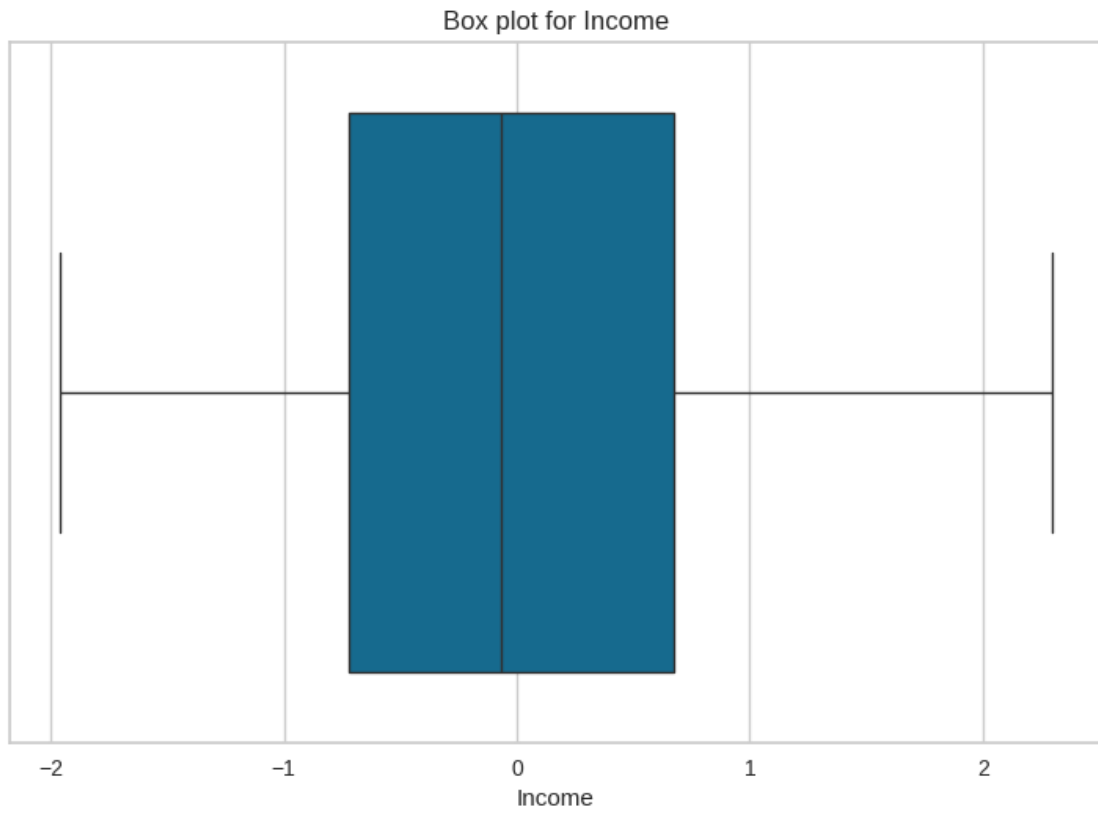
Series([], dtype: int64)




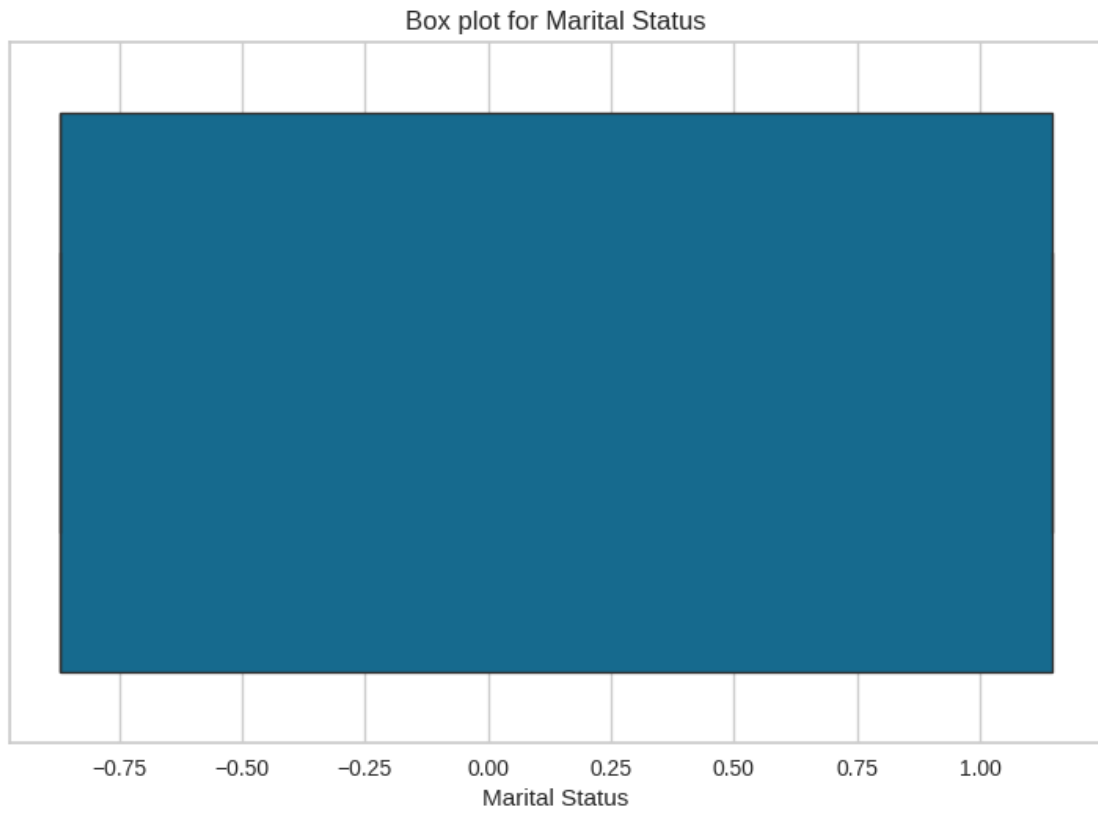
Number of Duplicates: 62






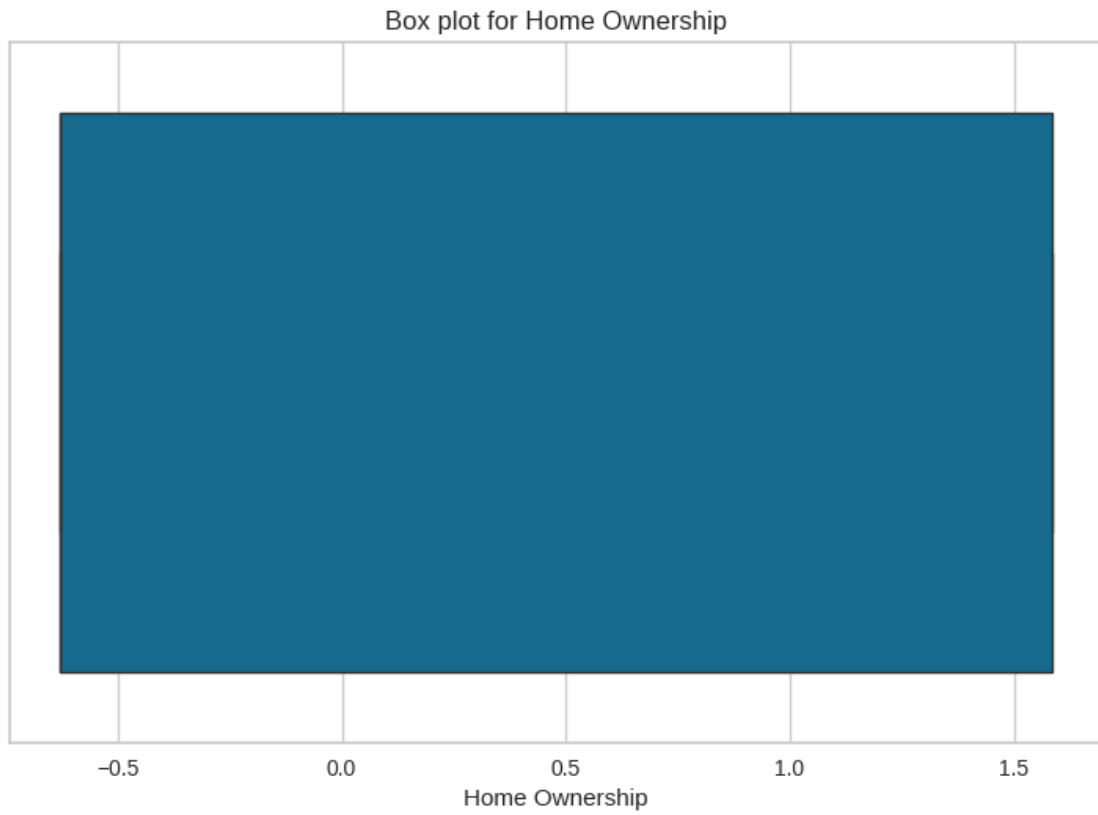
















Step 1: Checking for Missing Values

Missing Values in Each Column:

```
Series([], dtype: int64)
```




Step 2: Handling Missing Values

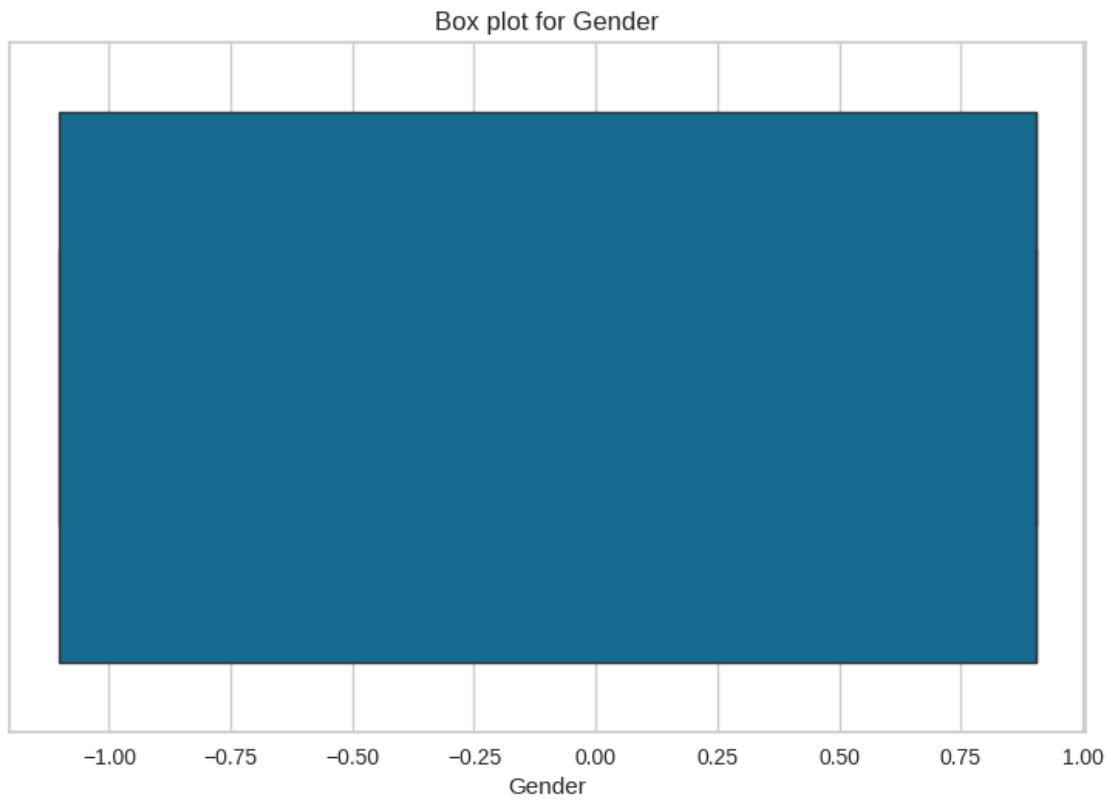
Step 3: Checking for Duplicates


Number of Duplicates: 0

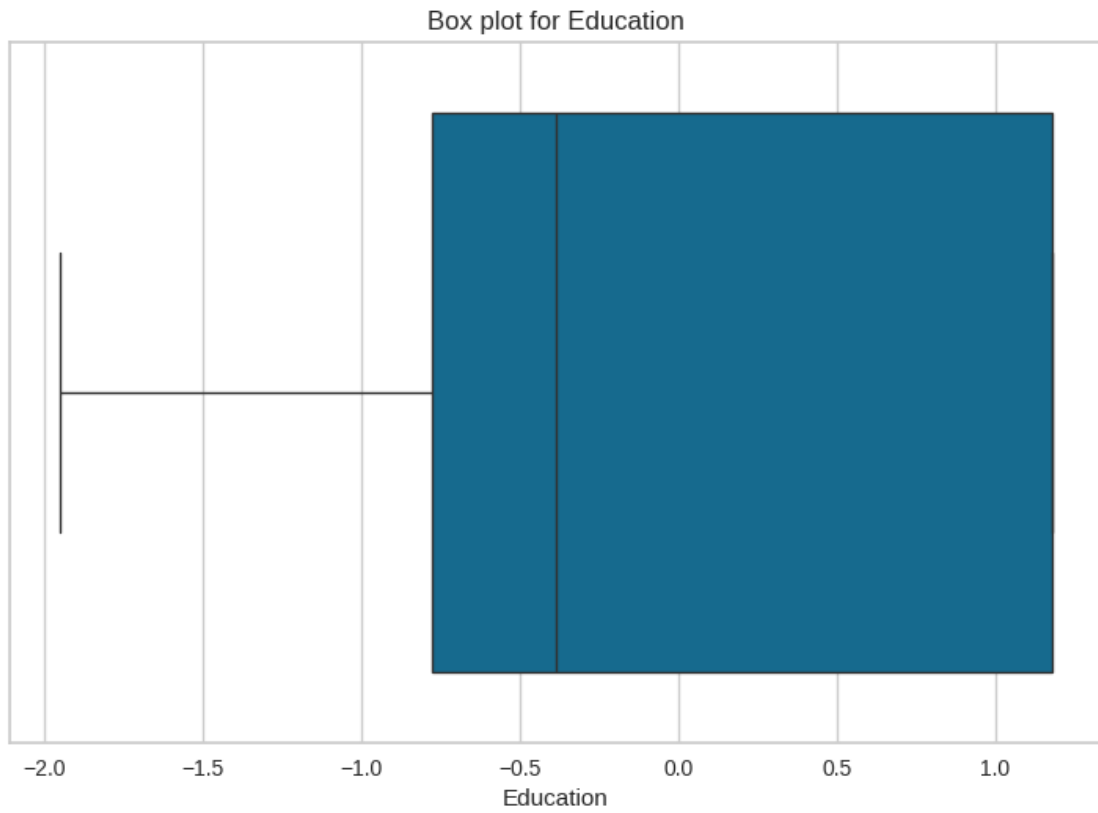
Step 4: Scaling Numeric Features

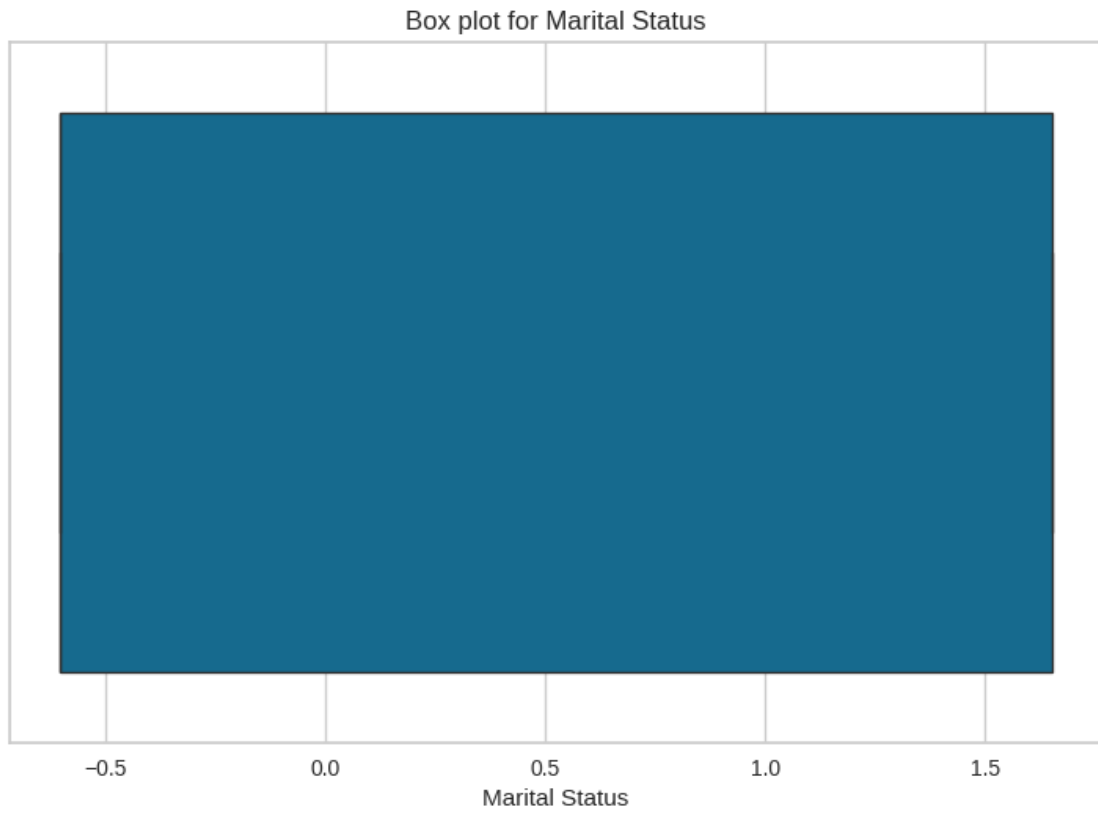
Step 5: Checking for Outliers



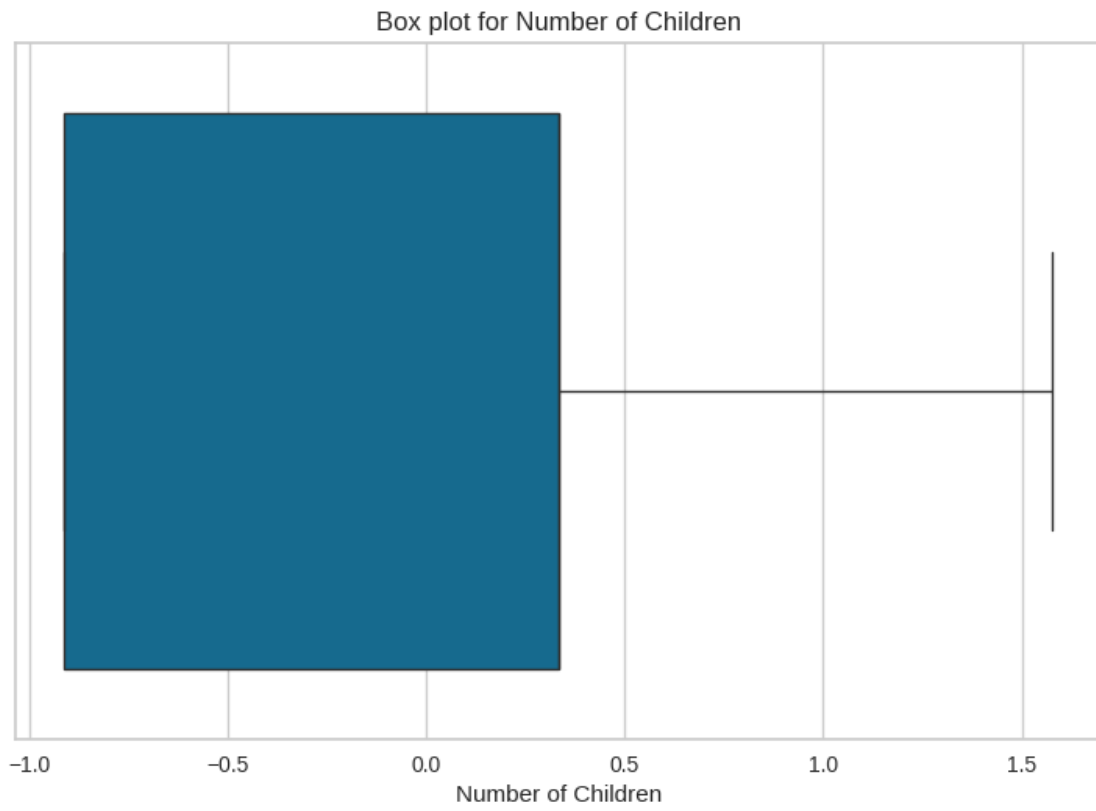


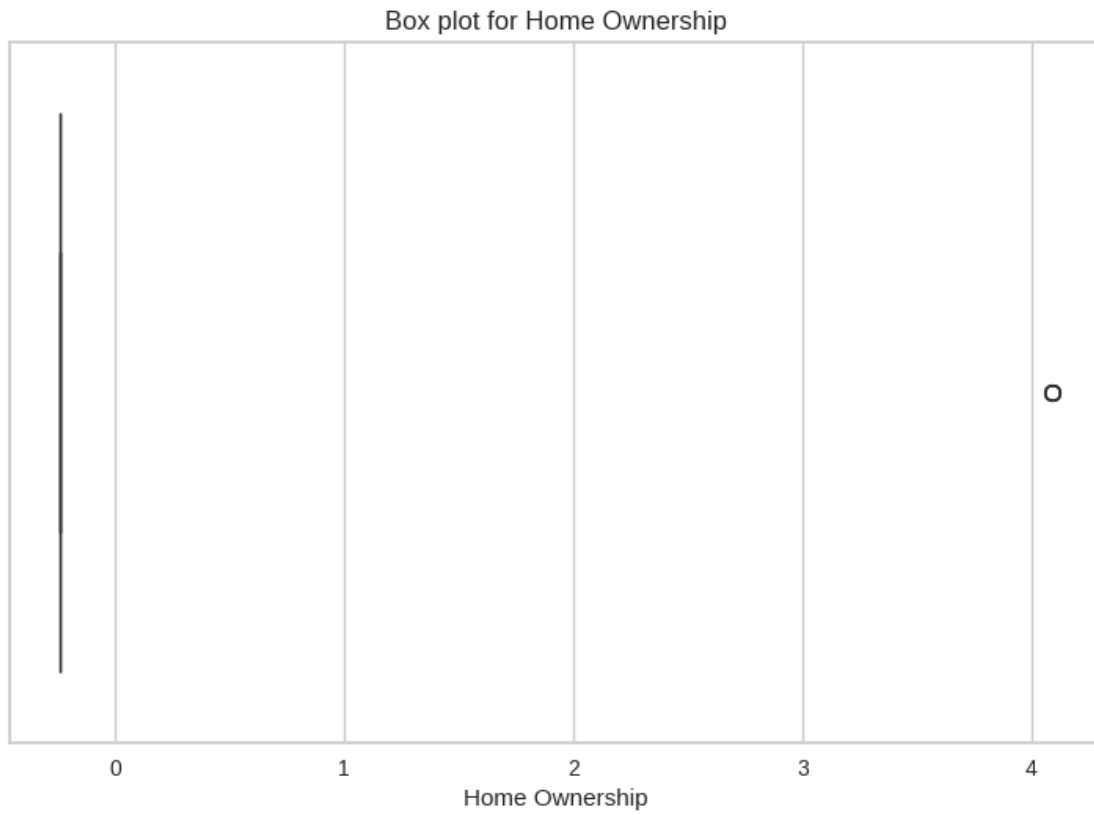


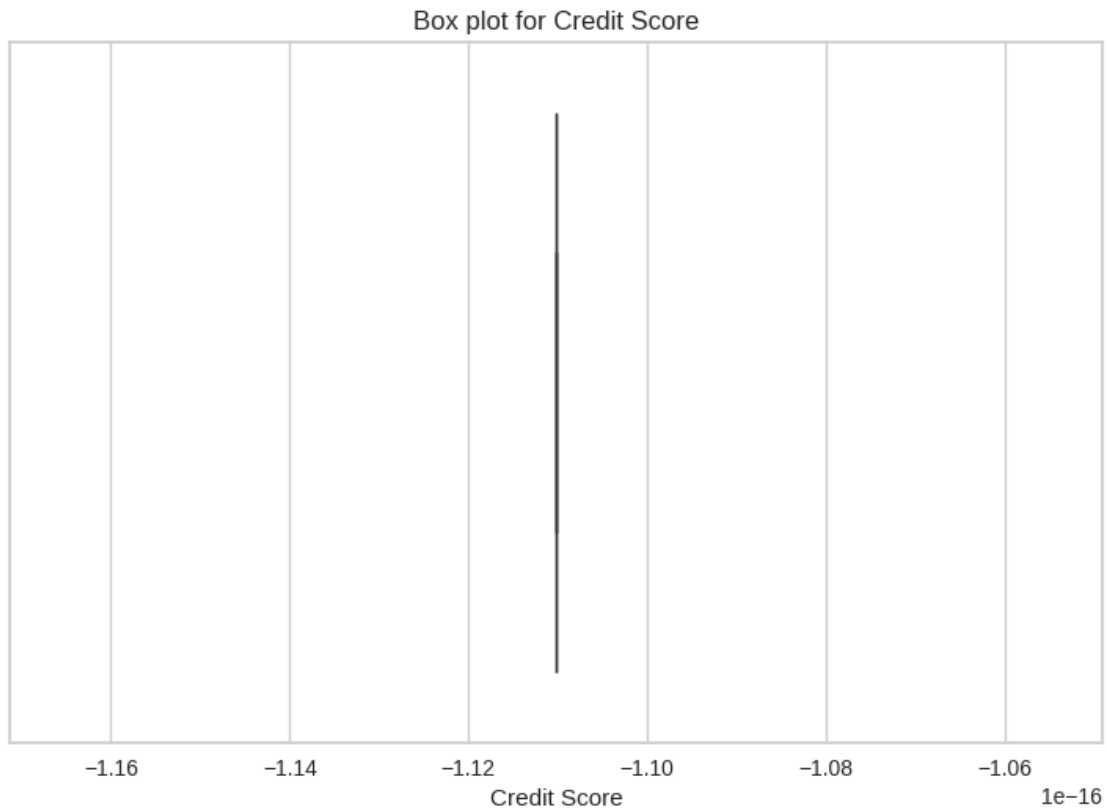












Step 6: Handling Outliers

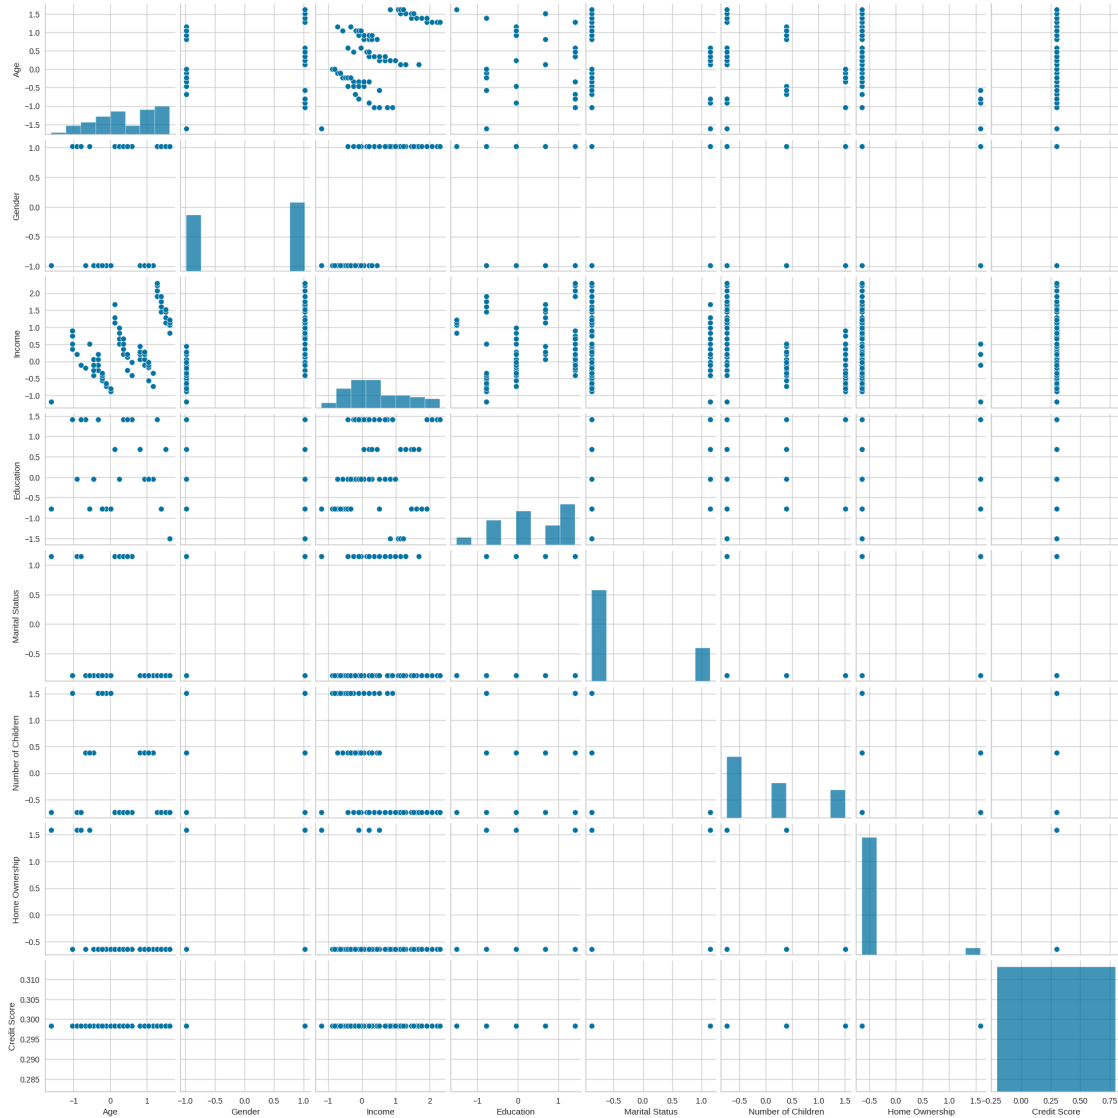
Data Cleaning Completed

## 11 Visualizations

### 12 1. Pair Plot


```
[35]: sns.pairplot(df)
```

```
[35]: <seaborn.axisgrid.PairGrid at 0x7f3d7de83250>
```



## 13 2. NULL Plot

```
[43]: import missingno as msno
      # Null count analysis
      null_plot = msno.bar(df)
```



## 14 3. Other Important Plots

```
[44]: def plots(df, variable):
    if df[variable].dtype != object:
        # define figure size
        fig, ax = plt.subplots(1, 5, figsize=(24, 4))

        # histogram
        sns.histplot(df[variable], bins=30, kde=True, ax=ax[0])
        ax[0].set_title('Histogram')

        # KDE plot
        sns.kdeplot(df[variable], ax=ax[1])
        ax[1].set_title('KDE Plot')






        # Line plot
        sns.lineplot(df[variable], ax=ax[2])
        ax[2].set_title('Line Plot')

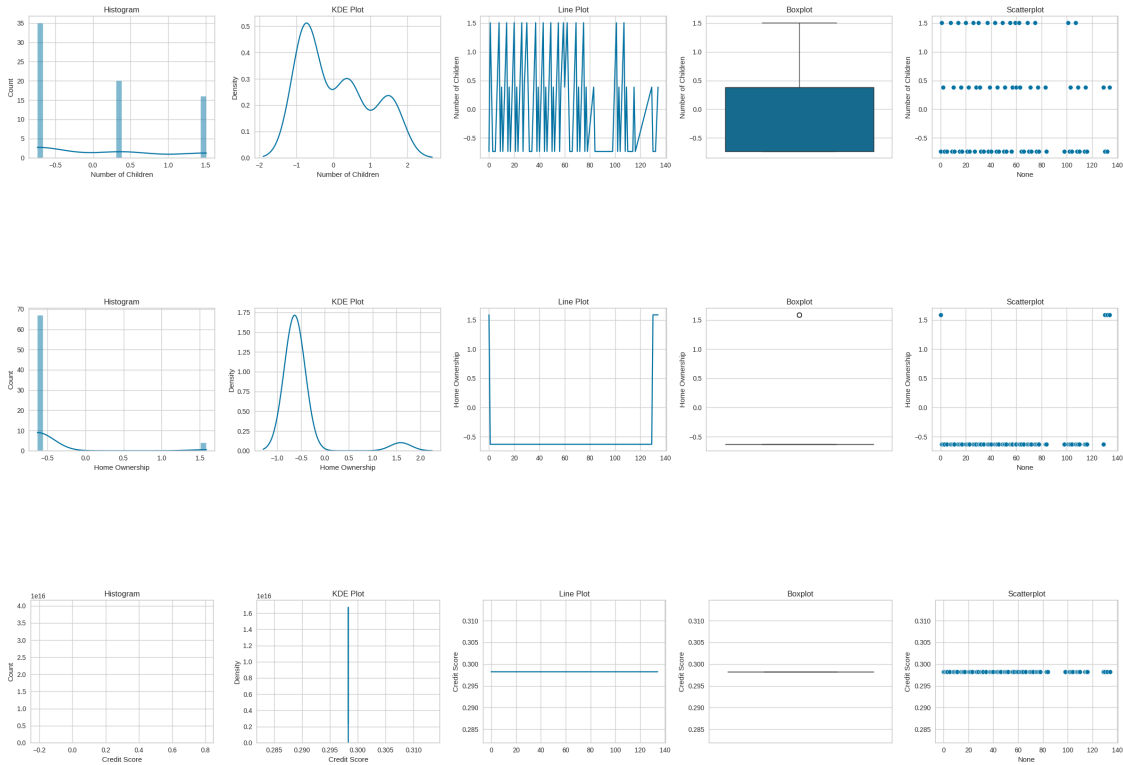
        # boxplot
        sns.boxplot(y=df[variable], ax=ax[3])
        ax[3].set_title('Boxplot')

        # scatterplot
        sns.scatterplot(x=df.index, y=df[variable], ax=ax[4])
        ax[4].set_title('Scatterplot')

    plt.tight_layout()
    plt.show()
```

```
for i in df.columns:
    plots(df ,i)
```





## 15 ML Modelling

```
[45]: # combine X_train_res and y_train_res
train_data = pd.concat([X_train_res, y_train_res], axis=1)

from pycaret.classification import *
s = setup(data=train_data, target='Credit Score', session_id=123,
normalize=True)
```

<pandas.io.formats.style.Styler at 0x7f3d6e1d0b50>

## 16 1. Comparing Models

```
[46]: compare_models()
```

<IPython.core.display.HTML object>

<pandas.io.formats.style.Styler at 0x7f3d7dca8f40>

Processing: 0% | 0/65 [00:00<?, ?it/s]

<IPython.core.display.HTML object>

```
[46]: KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
                           metric_params=None, n_jobs=-1, n_neighbors=5, p=2,
                           weights='uniform')
```

## 17 2. KNN Model

```
[47]: knn = create_model('knn')
```

```
<IPython.core.display.HTML object>
<pandas.io.formats.style.Styler at 0x7f3d7bf8c6a0>
Processing: 0%|          | 0/4 [00:00<?, ?it/s]
<IPython.core.display.HTML object>
```

```
[48]: preds = predict_model(knn)
```

```
<pandas.io.formats.style.Styler at 0x7f3d7a2ddcc0>
```

```
[49]: preds
```

```
[49]:
```

	Age	Gender	Income	Education	Marital Status	Number of Children \
140	29	0	47500	0	1	0
268	27	0	32785	0	1	0
228	27	0	37500	3	1	0
205	28	0	32500	0	1	0
103	43	1	92500	4	1	0
..	...	...	...	...	...	...
11	29	0	68000	2	0	2
246	28	0	32148	0	1	0
238	28	0	32037	0	1	0
85	27	0	37500	3	1	0
76	50	1	155000	4	0	0

	Home Ownership	Credit Score	prediction_label	prediction_score
140	1	0	0	1.0
268	1	2	2	1.0
228	1	2	2	1.0
205	1	2	2	1.0
103	0	1	1	1.0
..	...	...	...	...
11	0	0	1	1.0
246	1	2	2	1.0
238	1	2	2	1.0
85	1	2	2	1.0
76	0	1	1	1.0



[81 rows x 10 columns]

```
[50]: from sklearn.model_selection import cross_val_score

# Evaluate the ensemble model using cross-validation
scores = cross_val_score(knn, X_train_res, y_train_res, cv=20)
```


```
[51]: preds = predict_model(knn)
```

<pandas.io.formats.style.Styler at 0x7f3d7de55a50>

## 18 3. Plots


```
[52]: plot_model(knn)
```

<IPython.core.display.HTML object>




```
[53]: plot_model(knn, plot = 'error')
```

<IPython.core.display.HTML object>




```
[54]: plot_model(knn, plot = 'learning')
```

<IPython.core.display.HTML object>



```
[55]: plot_model(knn, plot = 'vc')
```

<IPython.core.display.HTML object>



## 19 4. Traditional Approach

```
[56]: from sklearn.neighbors import KNeighborsClassifier
```

```
[57]: clf = KNeighborsClassifier()

# Fit the Extra Trees Classifier object to the dataset
clf.fit(X_train_res, y_train_res)

scores = cross_val_score(knn, X_train_res, y_train_res, cv=20)

# Predict the labels for the test data
y_preds = clf.predict(X_test)
```

## 20 5. Classification Report

```
[58]: from sklearn.metrics import classification_report, confusion_matrix
print("Classification Report")
print(classification_report(y_test, y_preds))
```

Classification Report

	precision	recall	f1-score	support
0	1.00	1.00	1.00	5
1	1.00	1.00	1.00	23
2	1.00	1.00	1.00	5
accuracy			1.00	33
macro avg	1.00	1.00	1.00	33
weighted avg	1.00	1.00	1.00	33

## 21 100% Accuracy

## 22 6. Confusion Matrix

```
[59]: print("Confusion Matrix:")
print(confusion_matrix(y_test, y_preds))
```

Confusion Matrix:

```
[[ 5  0  0]
 [ 0 23  0]
 [ 0  0  5]]
```